


## STSCI4550Project

```
#####  
# State-Space Model for Air Data  
#####  
  
library(MARSS)  
load('~/.Downloads/Air.RData')  
#time(Air)  
  
# constructing a time plot of the data  
# days range from first 14 days in March 2004  
plot( x = time(Air)[1:336],  
      y = log( as.data.frame(Air)[1:336,"True CO"] ),  
      xlab='day',  
      ylab = "True CO",  
      main='Logged True CO Levels in 14 days in Mar 2004 in Italian City',  
      type = "l")
```


### Logged True CO Levels in 14 days in Mar 2004 in Italian City



```
# constructing a time plot of the data  
# days range from March 2004 until April 2005  
plot( x = time(Air),  
      y = log( as.data.frame(Air)[,"True CO"] ),  
      xlab='day', ylab = "Logged True CO",
```

```
main='Logged True CO Levels from Mar 2004 until April 2005 in Italian City',
type = "l")
```

## Logged True CO Levels from Mar 2004 until April 2005 in Italian City



```
# Create a 1 x 336 matrix containing hourly measurements of the
# log of True CO over the first 2 weeks (336/24=14):
```

```
ind <- 1:336
T <- length(ind)
y <- t(as.matrix(log(Air[ind,1])))
```

```
# Center/scale the exogenous variables:
```

```
temp <- scale(Air[ind, 3])
rhum <- scale(Air[ind, 4])
hum <- scale(Air[ind, 5])
```

```
### State space model specification in MARSS format:
```

```
# I'll provide some code which will help you to set up the
# state space model in MARSS format. You will have to fill
# in the rest.
```

```
# Constructing B (Phi in S&S notation):
```

```
# We first create a matrix of zeros (but we need to use a list
# because the matrix will not just contain numbers).
```

```
Blist <- list()
```

```

for(i in 1:(24^2)){
  Blist[[i]] <- 0
}
B <- matrix(Blist, nrow=24)
B[1,1] <- 'phi' # Top left entry is 'phi'
B[2,] <- c(0, rep(-1, 23)) # Filling in the second row
# We then fill in the remaining rows:
for(i in 3:24){
  B[i,i-1] <- 1
}

# You need to construct Q. You can do this in much the
# same way we constructed B.
Qlist <- list()
for(i in 1:(24^2)){
  Qlist[[i]] <- 0
}
Q <- matrix(Qlist, nrow=24)
Q[1,1] <- "q11" # Top left entry is 'phi'
Q[2,2] <- "q22"

# You need to construct x0 and V0
# (mu0 and Sigma0 in SES notation, respectively).
x0 <- as.matrix( rep(0,24) )
V0 <- diag(10,24)

# U has no equivalent in SES notation and the MARSS
# default value is not zero, so you need to set U equal
# to a zero vector of appropriate dimension.
U <- as.matrix( rep(0,24) )

# You need to construct the measurement matrix Z
# (A in SES notation).
Z <- rep(0,24)
Z[1:2] <- 1
Z <- t(Z)

# Constructing D (Gamma in SES notation):
D <- matrix(c('beta0', 'beta1', 'beta2', 'beta3'), nrow=1)

# Constructing d:

# We need to pass in the values of the exogenous variables at
# every point in time, so we use an array object as follows:

d <- array(rep(0, 4*length(ind)), dim=c(4,1,length(ind)))
for(t in ind){
  d[,t] <- c(1, temp[t], rhum[t], hum[t])
}

# You need to construct R (this is also R in SES notation).

```

```

R <- as.matrix( "r11" )

# A has no equivalent in SES notation and the MARSS
# default value is not zero, so you need to set A equal
# to a zero vector of appropriate dimension.
A <- as.matrix( 0 )

# Construct the model list as follows:
model.list <- list(B=B, Q=Q, x0=x0, V0=V0, U=U,
                  Z=Z, D=D, d=d, R=R, A=A, tinitx=0)

### Next you need to fit the model using the MARSS function.

# As we did in notes5code.R (e.g. with the biomarker data),
# you will need to first run the EM algorithm using the option
# method='kem' to get initial estimates and then pass these
# as initial values with method='BFGS'. During the first MARSS
# call (with method='kem') you should also pass in the argument
# control=list(minit=2, maxit=2) so that the EM algorithm only
# runs for a small number of iterations (even 2 iterations is fine).
fit_kem <- MARSS(y=y, model=model.list, control=list(minit=2, maxit=2), method='kem')

## Warning! Reached maxit before parameters converged. Maxit was 2.
## abstol not reached and no log-log test info since maxit less than min.iter.conv.test.
##
## MARSS fit is
## Estimation method: kem
## Convergence test: conv.test.slope.tol = 0.5, abstol = 0.001
## WARNING: Abstol convergence only no info on log-log convergence.
## No log-log convergence info because maxit (=2) < min.iter.conv.test (=15).
## The likelihood and params might not be at the ML values.
## Try setting control$maxit higher.
## Log-likelihood: -178.5309
##
##
##          Estimate
## R.r11      0.03672
## B.phi      0.96602
## Q.q11      0.03847
## Q.q22      0.02577
## D.beta0    0.00063
## D.beta1    0.00145
## D.beta2    0.00457
## D.beta3    0.00393
## Initial states (x0) defined at t=0
##
## Standard errors have not been calculated.
## Use MARSSparamCIs to compute CIs and bias estimates.
##
## Convergence warnings
## No convergence testing performed.

fit_bfgs_with_kem <- MARSS(y=y, model=model.list, method = "BFGS",
                           inits = fit_kem)

```

```

## Success! Converged in 60 iterations.
## Function MARSSkfas used for likelihood calculation.
##
## MARSS fit is
## Estimation method: BFGS
## Estimation converged in 60 iterations.
## Log-likelihood: -64.08381
## AIC: 144.1676   AICc: 144.6306
##
##           Estimate
## R.r11    9.15e-12
## B.phi    7.83e-01
## Q.q11    5.13e-02
## Q.q22    3.12e-13
## D.beta0  7.55e-01
## D.beta1  2.59e-02
## D.beta2  6.63e-02
## D.beta3  2.55e-02
## Initial states (x0) defined at t=0
##
## Standard errors have not been calculated.
## Use MARSSparamCIs to compute CIs and bias estimates.

```

```

# confidence intervals
MARSSparamCIs(fit_bfgs_with_kem)

```

```

##
## MARSS fit is
## Estimation method: BFGS
## Estimation converged in 60 iterations.
## Log-likelihood: -64.08381
## AIC: 144.1676   AICc: 144.6306
##
##           ML.Est  Std.Err  low.CI  up.CI
## R.r11    9.15e-12  0.005062 -0.009922 0.009922
## B.phi    7.83e-01  0.045043  0.694973 0.871537
## Q.q11    5.13e-02  0.008757  0.034089 0.068417
## Q.q22    3.12e-13  0.000159 -0.000312 0.000312
## D.beta0  7.55e-01  0.146444  0.467846 1.041897
## D.beta1  2.59e-02  0.140957 -0.250404 0.302136
## D.beta2  6.63e-02  0.135973 -0.200172 0.332831
## D.beta3  2.55e-02  0.063844 -0.099626 0.150639
## Initial states (x0) defined at t=0
##
## CIs calculated at alpha = 0.05 via method=hessian

```

```


# plot of state estimates
plot(fit_bfgs_with_kem)

```

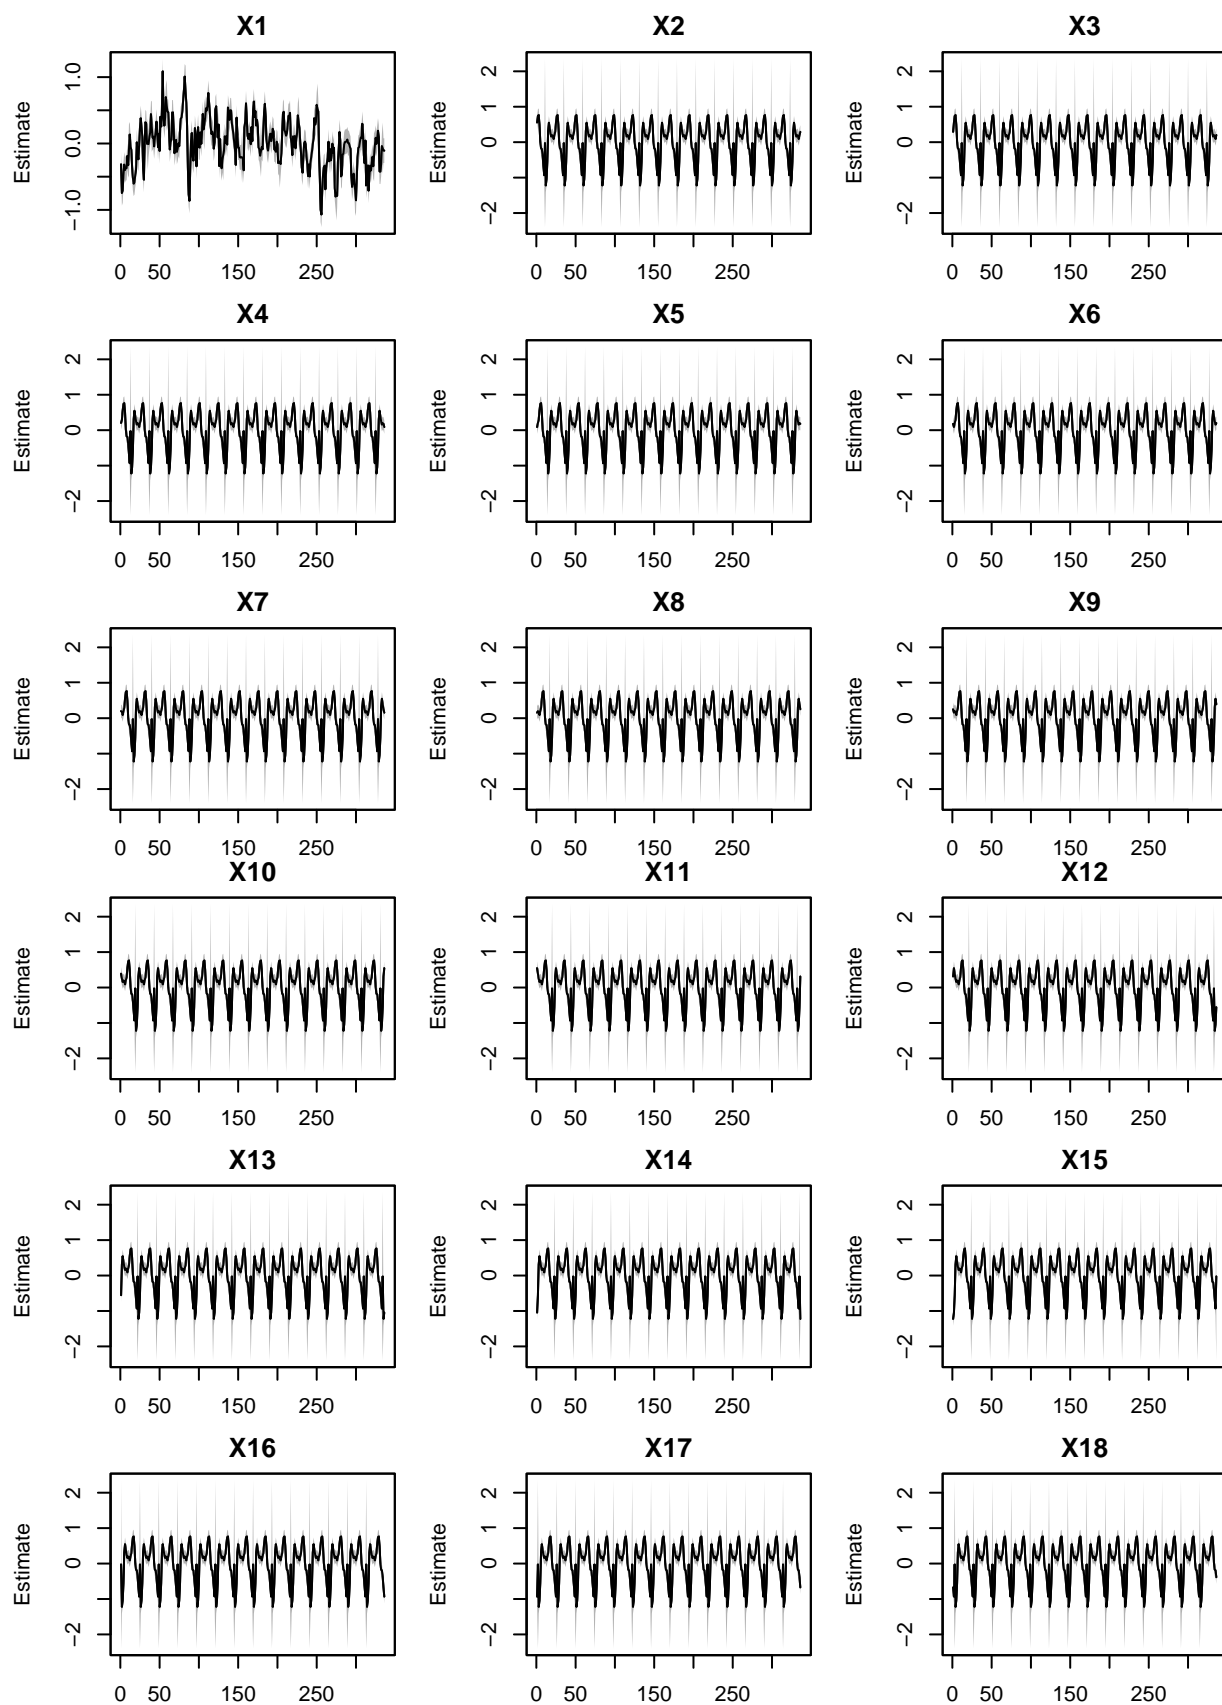
```


## MARSSresiduals.tt1 reported warnings. See msg element of returned residuals object.

```




```
## plot type = fitted.ytT Observations with fitted values  
## Hit <Return> to see next plot (q to exit):
```






```
## plot type = xtT Estimated states
## Hit <Return> to see next plot (q to exit):
```

**Y1**





```
## plot type = model.resids.ytt1
## Hit <Return> to see next plot (q to exit):
```






```
## plot type = std.model.resids.ytT  
## Hit <Return> to see next plot (q to exit):
```







```
## plot type = std.state.resids.xtT
## Hit <Return> to see next plot (q to exit):
```




```
## plot type = qqplot.std.model.resids.ytt1
## Hit <Return> to see next plot (q to exit):
```





```
## plot type = qqplot.std.state.resids.xtT
## Hit <Return> to see next plot (q to exit):
```



```
## plot type = acf.std.model.resids.ytt1
```

```
## FORECASTING
```

```
# Replace h values at the end of the y series (or the subset
# of the y series we're considering) with NAs which MARSS
# will then forecast:
```

```


h <- 24
ywithNAs <- y
ywithNAs[, (T-h+1):T] <- rep(NA, h)

fit_bfgs_with_kem <- MARSS(y=ywithNAs, model=model.list, method = "BFGS",
                           inits = fit_kem)

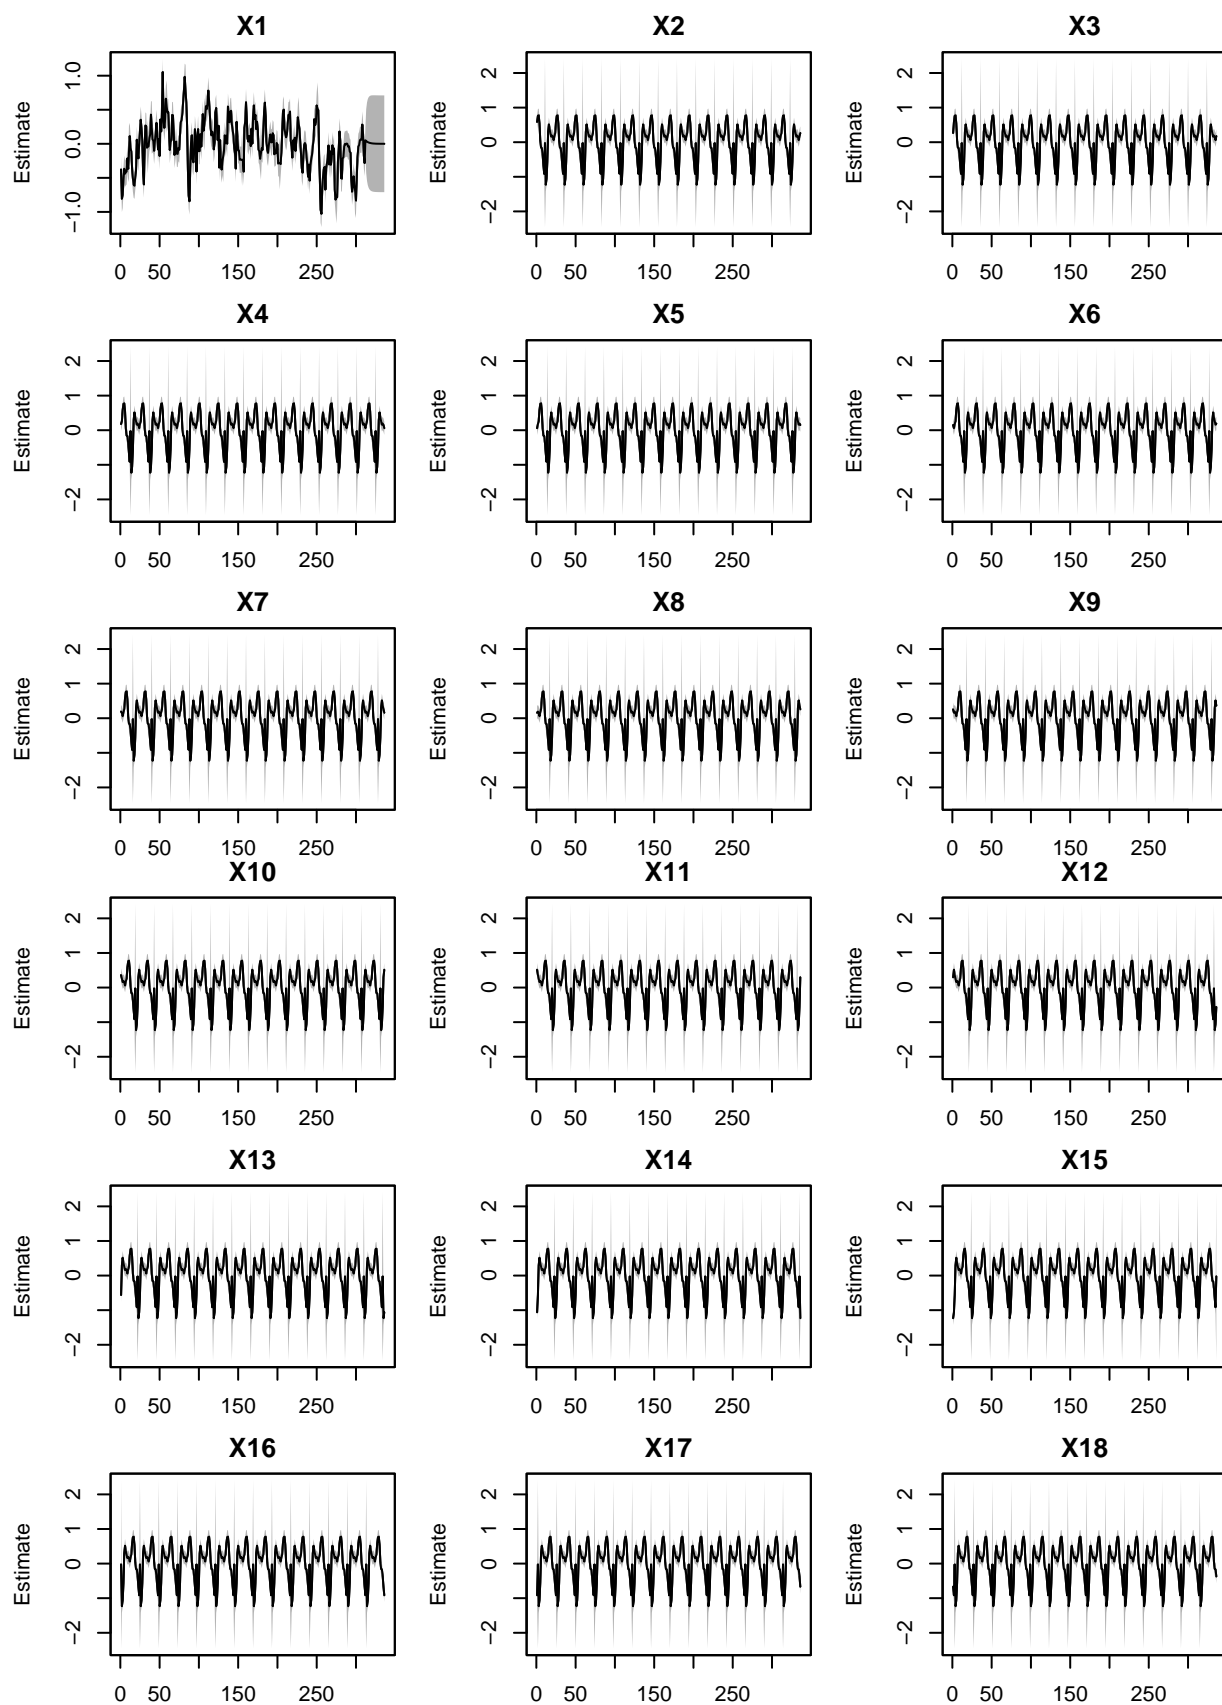
## Success! Converged in 131 iterations.
## Function MARSSkfas used for likelihood calculation.
##
## MARSS fit is
## Estimation method: BFGS
## Estimation converged in 131 iterations.
## Log-likelihood: -63.71363
## AIC: 143.4273   AICc: 143.9255
##
##           Estimate
## R.r11    2.18e-12
## B.phi    7.82e-01
## Q.q11    5.11e-02
## Q.q22    6.76e-11
## D.beta0  7.75e-01
## D.beta1  1.33e-02
## D.beta2  4.38e-02
## D.beta3  1.94e-02
## Initial states (x0) defined at t=0
##
## Standard errors have not been calculated.
## Use MARSSparamCIs to compute CIs and bias estimates.
plot(fit_bfgs_with_kem)

## MARSSresiduals.tt1 reported warnings. See msg element of returned residuals object.

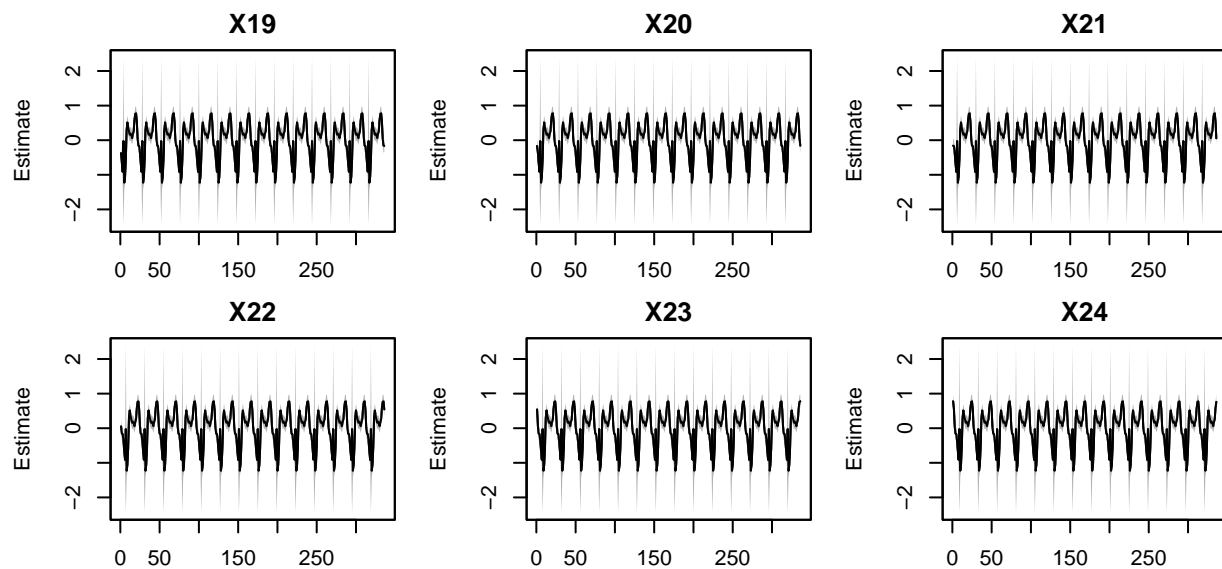
```



```
## plot type = fitted.ytT Observations with fitted values  
## Hit <Return> to see next plot (q to exit):
```









```
## plot type = xtT Estimated states
## Hit <Return> to see next plot (q to exit):
```


**Y1**




```
## plot type = model.resids.ytt1
## Hit <Return> to see next plot (q to exit):
```




```
## plot type = std.model.resids.ytT  
## Hit <Return> to see next plot (q to exit):
```







```
## plot type = std.state.resids.xtT
## Hit <Return> to see next plot (q to exit):
```




```
## plot type = qqplot.std.model.resids.ytt1
## Hit <Return> to see next plot (q to exit):
```





```
## plot type = qqplot.std.state.resids.xtT
## Hit <Return> to see next plot (q to exit):
```

**Y1**



```
## plot type = acf.std.model.resids.ytt1
y_forecasted <- fitted.values(fit_bfgs_with_kem)[(T-h+1):T, ".fitted"]
y_actual <- y[(T-h+1):T]
to_remove <- which(is.na(y_actual) == TRUE)

rmse <- sqrt(mean((y_actual[-to_remove] - y_forecasted[-to_remove])^2))
```

```

rmse

## [1] 0.3868232
sqrt( mean( ((y_actual - y_forecasted)^2), na.rm=TRUE) )

## [1] 0.3868232

#####
# Model 2: linear regression with ARMA errors#
#####

#####
# Initial analysis of the Air data #
#####

library(astsa) # For datasets and sarima function
library(forecast) # For Arima and forecast functions

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

##
## Attaching package: 'forecast'

## The following object is masked from 'package:astsa':
##
##   gas

load('~Downloads/Air.RData')

# Create a 1 x 336 matrix containing hourly measurements of the
# log of True CO over the first 2 weeks (336/24=14):

ind <- 1:336
T <- length(ind)
y <- t(as.matrix(log(Air[ind,1])))

# Center/scale the exogenous variables (except for time):

temp <- scale(Air[ind, 3])
rhum <- scale(Air[ind, 4])
hum <- scale(Air[ind, 5])
time <- time(Air)

# Split the data set into a training set, which we will use to
# fit the models, and a test set, which we will use to evaluate
# forecast accuracy.

h <- 48 # Size of test set

ytrain <- y[1:(T-h)]
temptrain <- temp[1:(T-h)]
rhumtrain <- rhum[1:(T-h)]
humtrain <- hum[1:(T-h)]
timetrain <- time[1:(T-h)]

```

```

ytest <- y[(T-h+1):T]
temptest <- temp[(T-h+1):T]
rhumtest <- rhum[(T-h+1):T]
humtest <- hum[(T-h+1):T]
timetest <- time[(T-h+1):T]

# Fit linear regression model to obtain residuals
# and then look at the ACF/PACF plot to determine ARMA order:


lmfit <- lm(ytrain ~ timetrain + temptrain + rhumtrain + humtrain, na.action=na.omit)
summary(lmfit)

##
## Call:
## lm(formula = ytrain ~ timetrain + temptrain + rhumtrain + humtrain,
##     na.action = na.omit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.83230 -0.30916 -0.03928  0.35073  1.44270
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.02009    0.08397  12.149 < 2e-16 ***
## timetrain    -0.04642    0.01335  -3.476 0.000593 ***
## temptrain    -0.28789    0.16090  -1.789 0.074703 .
## rhumtrain    -0.60382    0.17396  -3.471 0.000603 ***
## humtrain      0.23625    0.07690   3.072 0.002341 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5591 on 270 degrees of freedom
## (13 observations deleted due to missingness)
## Multiple R-squared:  0.1883, Adjusted R-squared:  0.1762
## F-statistic: 15.65 on 4 and 270 DF,  p-value: 1.561e-11

plot(lmfit$residuals)



```





```
acf2(lmfit$residuals)
```

### Series: Imfit\$residuals



```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
## ACF  0.77  0.41  0.09 -0.10 -0.16 -0.18 -0.19 -0.18 -0.08  0.06  0.19  0.18
## PACF 0.77 -0.46 -0.08  0.06 -0.02 -0.14 -0.04 -0.01  0.21  0.02  0.04 -0.23
##      [,13] [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24]
```

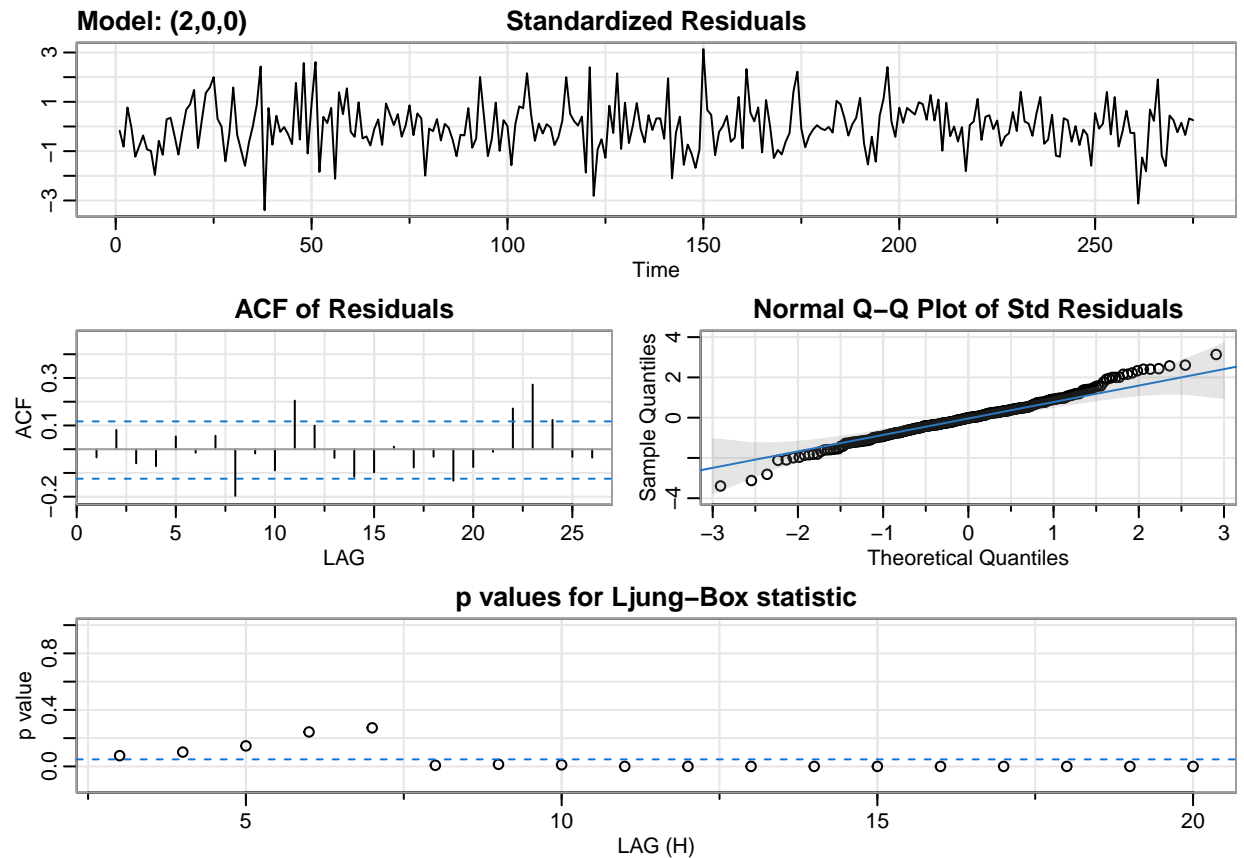
```
## ACF    0.05 -0.10 -0.20 -0.25 -0.29 -0.30 -0.25 -0.09  0.16  0.40  0.5  0.40
## PACF -0.06 -0.01 -0.03 -0.17 -0.12  0.02  0.09  0.14  0.17  0.14  0.0 -0.11
##      [,25] [,26] [,27]
## ACF    0.19  0.01 -0.11
## PACF -0.05  0.00 -0.06
```

```
# Fit an AR(2) model to the residuals.
```

```
# Diagnostic plots don't look amazing, but don't look horrible either:
```

```
sarima(lmfit$residuals, p=2, d=0, q=0)
```

```
## initial  value -0.587757
## iter    2 value -0.747315
## iter    3 value -1.010212
## iter    4 value -1.125561
## iter    5 value -1.140931
## iter    6 value -1.160254
## iter    7 value -1.160284
## iter    8 value -1.160286
## iter    9 value -1.160287
## iter   10 value -1.160288
## iter   11 value -1.160289
## iter   11 value -1.160289
## iter   11 value -1.160289
## final   value -1.160289
## converged
## initial  value -1.160167
## iter    2 value -1.160174
## iter    3 value -1.160175
## iter    4 value -1.160176
## iter    5 value -1.160177
## iter    6 value -1.160177
## iter    6 value -1.160177
## iter    6 value -1.160177
## final   value -1.160177
## converged
```



```
## $fit
##
## Call:
## arima(x = xdata, order = c(p, d, q), seasonal = list(order = c(P, D, Q), period = S),
##       xreg = xmean, include.mean = FALSE, transform.pars = trans, fixed = fixed,
##       optim.control = list(trace = trc, REPORT = 1, reltol = tol))
##
## Coefficients:
##          ar1          ar2      xmean
##          1.1227   -0.4546   0.0020
## s.e.    0.0535    0.0534   0.0567
##
## sigma^2 estimated as 0.09775:  log likelihood = -71.16,  aic = 150.32
##
## $degrees_of_freedom
## [1] 272
##
## $ttable
##      Estimate      SE t.value p.value
## ar1      1.1227 0.0535 20.9930 0.0000
## ar2     -0.4546 0.0534 -8.5059 0.0000
## xmean     0.0020 0.0567  0.0358 0.9715
##
## $AIC
## [1] 0.5466136
##
```

```
## $AICc
## [1] 0.5469356
##
## $BIC
## [1] 0.5992212

# Go ahead and fit the linear regression with AR(2) errors.
# The Arima/forecast functions are most convenient for subsequent forecasting:

fitAR2reg <- Arima(ytrain, order = c(2,0,0), xreg=cbind(time=timetrain, temp=temptrain, rhum=rhumtrain,


# Create data frames / data matrices to forecast the last h values of the series
h <- 48 # The number of steps ahead we want to forecast
T <- nrow(Air)

newdf <- data.frame(time=timetest, temp = temptest, rhum = rhumtest, hum = humtest)

newxreg <- model.matrix(~ -1 + time + temp + rhum + hum, data=newdf)

# Forecast plots
fc <- forecast(fitAR2reg , h=h, xreg = newxreg)
plot(fc)
```

## Forecasts from Regression with ARIMA(2,0,0) errors



```
# Calculate RMSE of forecasts:
sqrt(mean((fc$mean - ytest)^2, na.rm=TRUE))

## [1] 0.6719156
```